

Conduct Track: Make Legacy Move

Original text formatted for the UK AI Agent Hackathon EP5 track brief.

Price pool is 8,000 £

The problem

Large enterprises run on custom software built up over decades. When the business needs to change (a new pricing model, a new regulation, a new market) the software has to change too. But that software is millions of lines deep, with little documentation and few people who still understand how it works.

So a change that should take days takes months. Teams of consultants are hired, budgets run into the millions, and a few specialists become the bottleneck for the whole company. Closing this gap is what Conduct does. For this track, we want you to take it on, too.

Your mission

Pick a slow, inefficient process that happens at large enterprises today, and build a tool that lets a user do it far faster with AI, while staying in control. Start with the process. Use your knowledge of an industry, articles, or people who've worked in it to identify a specific task that enterprise teams spend weeks on. Then build something that does it in hours.

The key design choice: don't let the agent do everything on its own. In an enterprise, a fully autonomous agent is a non-starter. What works is keeping the user in control, with:

- A clear interface that shows what the agent is doing and why.
- Step-by-step execution that the user can follow, approve, or adjust.
- Real constraints on what the agent can change on its own.

The goal isn't to replace the expert. It's to make them ten times faster while they stay at the wheel.

What we're looking for

A real process. A task enterprise teams genuinely struggle with today, not one you invented for the demo. The user stays in control. Visibility into every step, the ability to work piece by piece, and clear limits on the agent. Control beats autonomy. A credible scale story. You won't have a real enterprise system to test on, and that's fine. Show you've thought about how this holds up on something far bigger and messier.

Getting started

To make your demo convincing, work with realistic data rather than something you made up. You don't need anything corporate: publicly available data works fine as stand-in "enterprise" data. A large open-source codebase on GitHub, a dataset from Kaggle, or any sizeable real-world data source all work well. The bigger and messier, the better the story.

No enterprise experience required. The hard part is picking the right process, which anyone can do.

Show us how AI lets a business change its systems as fast as it changes its mind.

Judging criteria

Technical execution (35%)	Does it actually work, reliably, on messy and ambiguous inputs, and recover from errors instead of only running in a scripted demo? We want real engineering, not a wrapper around a prompt.
Impact & speed-up (30%)	How dramatic is the improvement on a real, slow enterprise process? The strongest entries take something that takes a team weeks and make it happen in hours, with a clear sense of the time, cost, or expertise saved.
User stays in control (20%)	Does the user have visibility into every step, the ability to work through the process piece by piece, and clear limits on what the agent does on its own? A fully autonomous agent is a non-starter in an enterprise. Control beats autonomy.
The demo (20%)	Is there a clear, live before/after that a non-engineer can grasp in 90 seconds? A well-presented interface that makes the value obvious counts for a lot.